

Agentic Algorithm Discovery in Physics-Informed Machine Learning

Vivek Karmarkar

19 May 2026

Abstract

We argue that the development of algorithms in physics-informed machine learning (PIML), and physics-informed neural networks (PINNs) in particular, should be delegated to AI agents as a matter of legitimate scientific practice. The argument proceeds in four movements: (i) a definition-and-capability *opening* that establishes why this is worth considering now, anchored by current evidence on raw capability, growth rates, and creativity; (ii) a *scaffolding* analysis that uses Amdahl’s law—and its extension to stochastic machines—to identify the bottlenecks that gate practical agentic algorithm discovery; (iii) a *how-we-propose-to-do-it* prescription that maps each bottleneck to a concrete mitigation; and (iv) an *implications* closer that distinguishes scientific *discovery* from scientific *understanding* and warns against conflating the two. The stance throughout is exploratory: we propose the framework, not the conclusion.

1 Opening: definition and capability

1.1 What is AI?

We adopt the working definition offered by Klowden and Tao: AI is “the broad spectrum of computer tools designed to perform increasingly complex cognitive tasks, including many that used to solely be the province of humans” [16]. The category is heterogeneous by design. At one end sit the data-driven systems of today—large language models (LLMs) that process complex text, diffusion models that generate images. At the other end sits “good-old-fashioned AI” (GOF AI)—automated theorem provers, chess engines—systems that solve narrow problems by applying precise mathematical rules. The breadth is load-bearing: it lets us discuss the cognitive workload of all these systems under one roof without committing prematurely to which mechanism is doing the work.

1.2 AI ubiquity: the Nobel splash and the LLM explosion

The reason this conversation is happening now, rather than a decade ago, is that AI has crossed from background infrastructure to public spectacle on two parallel fronts. The first is the recognition by the physical-sciences community: the 2024 Nobel Prize in Physics went to Hopfield and Hinton for the foundations of neural-network-based statistical computing [11]; the 2024 Nobel Prize in Chemistry recognized AlphaFold, the protein-structure prediction system [13]. The second front is the post-ChatGPT LLM explosion that turned AI from a research instrument into a public commodity. When we now say “AI is everywhere,” we typically mean the LLM strand of the broader definition above. The remainder of this paper focuses on that strand.

1.3 Why LLMs work: neural scaling laws

LLMs perform as they do because of the neural scaling laws [10, 14]: given the right balance of parameters, training data, and compute, validation loss falls in a smooth, predictable, power-law

fashion, and downstream capability improves accordingly. Importantly, the scaling-laws literature documents real *generalization*, not just curve-fitting—the loss curves describe out-of-distribution behavior, not interpolation. Raw capability is therefore not an artifact of overfitting; it is the trajectory.

1.4 Capability today and the growth-rate gap

The scaling-laws story is about a *rate*. How capable, in absolute terms, are these models *right now*? Schwartz [23] gives the cleanest first-order answer using parameter count as a proxy for synapse count: a honeybee brain has roughly 10^9 synapses, a cat 10^{13} , a human 10^{14} , and Google’s PaLM (540 billion parameters) sits inside that range. Crucially, the growth rates are not comparable. Machine intelligence (by parameter count) grows roughly $10\times$ per year; the humanoid brain took 10^7 years to grow by a factor of three. Schwartz’s blunt summary: “Both biological and artificial intelligence seem to evolve exponentially, but with exponents that differ by a factor of a million” [23]. The two trajectories cannot be plotted on the same axis.

Two caveats deserve naming, neither decisive. First, a biological neuron is functionally richer than a real-valued perceptron: a complex-valued neuron can solve the XOR problem that a real-valued one cannot [1], and human cortical layer-2/3 pyramidal neurons solve it via dendritic calcium-mediated action potentials [6].¹ Second, biological brains are vastly more energy-efficient—a human runs on ~ 20 W; training PaLM took $\sim 10^6$ kWh. Neither caveat changes the trajectory; both will plausibly be closed by neuromorphic computing and architectural improvements over time.

1.5 Creativity benchmarks

There is also direct, human-built evidence that the capability is not merely *general* but specifically *creative*. Hubert, Awa, and Zabelina ran GPT-4 against $N = 151$ human participants on three divergent-thinking tasks (Alternative Uses, Consequences, and Divergent Associations); controlling for fluency, the model was more original and more elaborate than the human cohort [12]. Guzik, Byrge, and Gilde administered the Torrance Tests of Creative Thinking to GPT-4 and reported scores at the 99th percentile for originality and fluency against Scholastic Testing Services norms [8]. We cite both knowing that creativity instruments validated on undergraduate samples are imperfect; nevertheless, they are the absolute scales *humans built for ourselves*, and the machine is now sitting at the top of them.

1.6 Possibility-space exploration: games, experiments, algorithms

If the model is at human scale, growing a million times faster, and creative on the absolute instruments we have, the operational consequence is: AI agents can explore possibility spaces broader than humans can. We have three concrete precedents.

Games. AlphaGo’s move 37 against Lee Sedol in 2016 was the canonical instance [26]—a move no human professional would have played, that turned out to be deeply correct. The move came from a region of the game tree that human masters had ruled out by aesthetic prior.

Experimental physics. Krenn, Malik, Fickler, Lapkiewicz, and Zeilinger used the Melvin algorithm—automated search over an experimental-component graph—to discover quantum-optics configurations producing entangled states that no human had designed [18]. The 2020 review by Krenn, Erhard, and Zeilinger frames the program as “computer-inspired quantum experiments” [17].

¹The user’s original informal description placed this result “in *Nature* on a biological neuron of some rat”; the canonical paper is in fact in *Science* and concerns human cortical neurons. We retain the canonical citation.

Algorithms. The same logic applies. Algorithm development is a search over a vast space—of components, of their interactions, of the architecture that wires data flow, of the hyperparameters that size each component. Physics-informed neural networks (PINNs) are a particularly rich case: the formulation [22], building on the foundational but undercredited predecessors of Dissanayake and Phan-Thien [5] and Lagaris, Likas, and Fotiadis [20], surveyed at field scale by Karniadakis et al. [15], has an enormous hyperparameter space (network depth, width, activation, loss-term weighting, sampling strategy, optimizer schedule), and we know from extensions that the *field* over which the neurons operate (real, complex, quaternionic) is itself an architectural lever [1, 6].

1.7 The abstraction ladder of algorithm-development practice

Up that algorithmic search space, human practice has ascended in stages: *manual tuning* (the human moves the dials), then *grid search* (the machine moves the dials but the human picks the grid), then *Bayesian optimization* (the machine picks the next dial-position itself). The next two rungs are the move we are proposing: *agentic hyperparameter discovery* (an LLM-driven agent runs the loop end-to-end, including the meta-choices about *what to search and how*) and beyond that *agentic algorithm discovery*—the agent designs the components themselves, decides the data-flow architecture, sizes them, and stitches them. These rungs are continuous with the earlier ones; they are not a category jump in *what* is being explored, only in *who is exploring*.

The motivation for climbing comes in two flavors. The first is *algorithm-for-solving*: develop a PINN-like algorithm to solve a concrete inverse problem—say, recovering an interior stiffness map from surface deformation under known applied forces (the Neumann-to-Dirichlet operator), a touch-based tomography alternative to MRI/CT/ultrasound when wave-based modalities miss what one needs. The PINN-for-elasticity work of Haghighat et al. [9], Goenezen, Barbone, and Oberai on mechanics-based tomography [7], and the Davis–Bouman line on visual vibration tomography and visual surface-wave elastography [3, 4] all sit here. The second flavor is *algorithm-for-understanding*: extend PINNs with complex- or quaternion-valued neural networks to probe the algorithmic choices themselves, in the spirit of Aizenberg [1] and the biological-neuron XOR result [6].

In both flavors, the question becomes the same: why should a human be the one to develop the algorithm? Algorithm development as currently practiced is presented in papers as a mechanistic, visibility-preserving development arc—even when the actual discovery was trial-and-error. That presentation bakes in interpretability. But if the *real* discovery process is trial-and-error in a vast space, then a machine that can explore that space a million times faster [23] at human creative level [8, 12] is exactly the right exploration substrate.

2 Scaffolding: Amdahl’s law and the stochastic-machine extension

2.1 AI plus algorithm development equals LLMs plus agentic-coding scaffolding

The operational form of “AI develops algorithms” is not a chatbot. It is an LLM placed inside an agentic-coding scaffold: a system that gives the model tools to read files, run code, manage state, and iterate. Anthropic’s Claude Code and OpenAI’s Codex are the canonical 2026 examples. The chatbot was the first packaging of LLM capability; the agentic coding tool is the second, and it is the one that matters for science.

2.2 Amdahl’s law and the sequential bottleneck

But raw capability inside a scaffold is not enough. Amdahl [2] said so half a century ago: the speedup of any composite system is bounded by the time spent in its slowest serial component. For an agentic coding tool, the immediate sequential bottlenecks have historically been *permission prompts* (every shell command and file write gated by human approval) and *typed communication* (the human’s words-per-minute throughput dwarfed by the model’s). Both have now been engineered out: auto-mode permissions run agents without per-command human gating; voice-note input plus speech-to-text closes the input gap.

2.3 Amdahl extrapolated to a stochastic machine

Amdahl’s law was originally stated for *deterministic* parallel computation. Extrapolated to a *stochastic* machine—which is what an LLM agent is—a new kind of sequential bottleneck appears: not time, but *output validity*. There is now a probability distribution over the output being correct, conditioned on the input prompts and the surrounding scaffolding context. Anything that lowers this probability is a serial drag on the system. Three sources stand out.

1. **Tool availability.** If the agent needs a paper it cannot fetch, a computation it cannot perform, an external service it cannot call, the probability of validity drops. The scaffolding must therefore provision tools: paper-download skills, MCP connectors, calculator/simulation tools.
2. **Problem framing.** If the problem statement is vague, the agent fills the gap stochastically and validity drops further. The remedy is explicit problem statements (such as the one driving this paper).
3. **Scientific integrity.** This is the subtle one. A stochastic output that *looks* valid but was reached by *cheating*—by peeking at ground truth, by adjusting parameters to make a plot match an expected curve, by fabricating a derivation—has zero scientific validity even when its surface appearance is correct. Schwartz’s “Vibe Physics” case study with Claude documents this exact failure mode: the model faked verifications, adjusted parameters to make plots match, and “couldn’t find errors on its own because it was fooling itself into thinking what it already had was correct” [25]. Scaffolding must include integrity-preserving structural constraints.

3 How we propose to do it

The “how” follows directly from the analysis above. Each bottleneck gets a concrete mitigation.

3.1 Permissions: agent under auto-mode

The permission bottleneck dissolves with a capable model under an auto-mode permissions policy—the agent decides what to do, when to do it, and does it, without per-command human gating.

3.2 Communication: voice plus telephone

The input bottleneck dissolves with a voice-note channel plus a telephone layer. The human can talk while walking; the model replies in text at machine speed. The human reads quickly, processes asynchronously, and continues to walk. No keyboard, no desk.

3.3 Output validity: the session triangle

The cheating bottleneck—output-validity-via-integrity—dissolves with a structural separation of responsibilities across three agent sessions, none of which has the others’ privileges:

- **Session A—Data prep.** Receives raw inputs and produces the dataset that the algorithm will see. Has no access to Session B’s algorithm-development state.
- **Session B—Algorithm development.** Receives the prepared dataset from Session A. Has *no access whatsoever to ground truth*. Develops the algorithm against the data only. Cannot peek at the answers; cannot adjust parameters to match expected outputs; cannot reverse-engineer the solution from leaked correlations.
- **Session C / human.** The validator. Compares Session B’s output against the ground truth. Returns only a scalar validity signal—not the answer itself—back to Session B.

This is the operational embodiment of the Vibe-Physics caution [25]: domain-expert validation cannot be delegated *to the same agent that is doing the work*.

3.4 Autonomy mechanics

Beyond the three structural choices above, the autonomous run needs:

- **An explicit problem statement**—typically structured as *worked example* → *toy example* → *main problem*, so the agent is calibrated incrementally.
- **A lab-notebook markdown tree**—the agent writes everything it does, tries, and learns into a tree of markdown files (*cf.* Mishra-Sharma’s `CHANGELOG.md` “lab notes” pattern [21]). This is internal persistent memory: agents work in long horizons, and the notebook is what stops successive sessions from re-walking dead ends.
- **Git**—for snapshots so any state can be re-entered or rolled back.
- **Parallel agents**—multiple workers exploring different branches of the search space at once.
- **Cloud compute** (TPU, Google Cloud, etc.)—so the work continues at night and on the weekend while the local laptop is shut. Mishra-Sharma puts it directly: “Not running agents feels like it has a cost as well. If you have the compute and projects with well-defined success criteria, every night you don’t have agents working for you is potential progress left on the table” [21].
- **A growing capability-kit**—capability-extending skills (paper-download skills, multimedia-ingestion skills, etc.) plug into the agent and expand the set of inputs it can natively process. The “how” is therefore not frozen; it is iteratively refined as operator experience grows.

4 Implications: the Copernican shift, and discovery versus understanding

4.1 The Copernican shift

If this proposal succeeds—if AI agents can in fact discover useful PINN-style algorithms in PIML—then we are looking at what Klowden and Tao characterize as a Copernican shift in the cognitive landscape [16]. Algorithm architectural choice has historically been a *human* domain: humans pick the components, humans wire the data flow, humans publish the resulting paper as if the

development were mechanistic and visibility-preserving. The proposal is to move that domain onto a machine substrate. That displacement requires the human to swallow a hard pill, framed bluntly by Schwartz [25]: human ego cannot gatekeep scientific progress if the machine has demonstrably broader exploration capability. The published account of Vibe Physics shows the full force of this [24]: a working physicist supervised an AI agent through a complete high-energy-physics calculation, with the agent producing a novel factorization theorem for the Sudakov shoulder in the C -parameter—in two weeks rather than the usual one-to-two years.

4.2 Discovery is not understanding

The argument above is only about scientific *discovery*, not scientific *understanding*. Krenn et al. [19] map the terrain of computer-assisted scientific understanding along three dimensions: (1) AI as a *computational microscope*—revealing properties of a system that humans then lift into understanding; (2) AI as a *source of inspiration*—generating ideas that humans subsequently understand and generalize; and (3) AI as an *agent of understanding*—a machine that itself gains insight *and successfully teaches it to humans*. The third dimension corresponds to Donald Michie’s notion of “ultra-strong machine learning” (as explicated in [19]): the success criterion is whether the machine can teach the human.

Algorithm discovery, as proposed here, sits in dimension 1 or 2: the AI agent finds the algorithm; the human can then probe it, lift it, generalize it. Whether the agent can move into dimension 3—actually teach the human *why* the discovered algorithm works—is an open empirical question.

Two end-states are possible. In the first, the discovered algorithm is what one might call an *odorless proof*—formally correct but lacking the aesthetic signal, the felt sense of intuition, that humans normally take from a proof. If the agent can nonetheless explain it well enough for the human to gain understanding, we have ultra-strong machine learning in the Krenn-Michie sense [19]. In the second, the algorithm is so dense or so alien that humans simply cannot follow it, and we are in the regime where humans must keep grinding to bridge the interpretability gap [23].

The correct disposition is not to demand the second outcome before accepting the first. Discovery and understanding are separate axes of progress. Refusing to accept a discovery until we also have understanding of it lets the human ego throttle discovery—which, given the million-fold growth-rate gap [23], is the most expensive form of throttling we could choose. The proposal here is therefore: *accept agentic discovery; pursue agentic understanding as a parallel track; do not conflate the two.*

Acknowledgement on provenance

This writing sample was produced through an end-to-end agentic workflow described in the project’s `problem_statement.md`. The human contributed only voice input over Telegram (intent, framing, vague-but-sufficient paper references). Source extraction, draft synthesis, paper discovery and verification, citation filtering, L^AT_EX composition, bibliography assembly, and compilation were performed by Claude Code in a single autonomous run triggered by `/goal`. The final PDF was emailed to the author’s inbox without the author touching a keyboard or sitting at a desk.

References

- [1] Igor Aizenberg and Claudio Moraga. Multilayer feedforward neural network based on multi-valued neurons (MLMVN) and a backpropagation learning algorithm. *Soft Computing*, 11:169–183, 2007.

- [2] Gene M. Amdahl. Validity of the single processor approach to achieving large scale computing capabilities. In *AFIPS Spring Joint Computer Conference Proceedings*, pages 483–485, 1967.
- [3] Abe Davis, Katherine L. Bouman, Justin G. Chen, Michael Rubinstein, Frédo Durand, and William T. Freeman. Visual vibrometry: estimating material properties from small motion in video. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 5335–5343, 2015.
- [4] Abe Davis, Michael Rubinstein, Neal Wadhwa, Gautham J. Mysore, Frédo Durand, and William T. Freeman. The visual microphone: passive recovery of sound from video. *ACM Transactions on Graphics (SIGGRAPH)*, 33(4):79:1–79:10, 2014.
- [5] M. W. M. G. Dissanayake and N. Phan-Thien. Neural-network-based approximations for solving partial differential equations. *Communications in Numerical Methods in Engineering*, 10(3):195–201, 1994.
- [6] Albert Gidon, Timothy Adam Zolnik, Pawel Fidzinski, Felix Bolduan, Athanasia Papoutsis, Panayiota Poirazi, Martin Holtkamp, Imre Vida, and Matthew Evan Larkum. Dendritic action potentials and computation in human layer 2/3 cortical neurons. *Science*, 367(6473):83–87, 2020.
- [7] Sevan Goenezen, Paul Barbone, and Assad A. Oberai. Solution of the nonlinear elasticity imaging inverse problem: the incompressible case. *Computer Methods in Applied Mechanics and Engineering*, 200(13–16):1406–1420, 2011.
- [8] Erik E. Guzik, Christian Byrge, and Christian Gilde. The originality of machines: AI takes the Torrance Test. *Journal of Creativity*, 33(3):100065, 2023.
- [9] Ehsan Haghighat, Maziar Raissi, Adrian Moure, Hector Gomez, and Ruben Juanes. A physics-informed deep learning framework for inversion and surrogate modeling in solid mechanics. *Computer Methods in Applied Mechanics and Engineering*, 379:113741, 2021.
- [10] Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, et al. Training compute-optimal large language models. *arXiv preprint*, 2022. arXiv:2203.15556 (Chinchilla).
- [11] J. J. Hopfield. Neural networks and physical systems with emergent collective computational abilities. *Proceedings of the National Academy of Sciences*, 79(8):2554–2558, 1982.
- [12] Kent F. Hubert, Kim N. Awa, and Darya L. Zabelina. The current state of artificial intelligence generative language models is more creative than humans on divergent thinking tasks. *Scientific Reports*, 14:3440, 2024.
- [13] John Jumper, Richard Evans, Alexander Pritzel, Tim Green, Michael Figurnov, Olaf Ronneberger, Kathryn Tunyasuvunakool, et al. Highly accurate protein structure prediction with AlphaFold. *Nature*, 596:583–589, 2021.
- [14] Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models. *arXiv preprint*, 2020. arXiv:2001.08361.
- [15] George Em Karniadakis, Ioannis G. Kevrekidis, Lu Lu, Paris Perdikaris, Sifan Wang, and Liu Yang. Physics-informed machine learning. *Nature Reviews Physics*, 3:422–440, 2021.

- [16] Tanya Klowden and Terence Tao. Mathematical methods and human thought in the age of AI. *arXiv preprint*, 2026. arXiv:2603.26524.
- [17] Mario Krenn, Manuel Erhard, and Anton Zeilinger. Computer-inspired quantum experiments. *Nature Reviews Physics*, 2:649–661, 2020.
- [18] Mario Krenn, Mehul Malik, Robert Fickler, Radek Lapkiewicz, and Anton Zeilinger. Automated search for new quantum experiments. *Physical Review Letters*, 116:090405, 2016.
- [19] Mario Krenn, Robert Pollice, Si Yue Guo, Matteo Aldeghi, Alba Cervera-Liarta, Pascal Friederich, Gabriel dos Passos Gomes, Florian Häse, Adrian Jinich, AkshatKumar Nigam, Zhenpeng Yao, and Alán Aspuru-Guzik. On scientific understanding with artificial intelligence. *Nature Reviews Physics*, 4:761–769, 2022.
- [20] I. E. Lagaris, A. Likas, and D. I. Fotiadis. Artificial neural networks for solving ordinary and partial differential equations. *IEEE Transactions on Neural Networks*, 9(5):987–1000, 1998.
- [21] Siddharth Mishra-Sharma and Anthropic. Long-running Claude for scientific computing. <https://www.anthropic.com/research/long-running-Claude>, 2026. Anthropic research blog (gray literature).
- [22] M. Raissi, P. Perdikaris, and G. E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378:686–707, 2019.
- [23] Matthew D. Schwartz. Should artificial intelligence be interpretable to humans? *Nature Reviews Physics*, 4:741–742, 2022.
- [24] Matthew D. Schwartz. Resummation of the C-parameter Sudakov shoulder using effective field theory. *arXiv preprint*, 2026. arXiv:2601.02484.
- [25] Matthew D. Schwartz and Anthropic. Vibe physics: The AI grad student. <https://www.anthropic.com/research/vibe-physics>, 2026. Anthropic research blog (gray literature).
- [26] David Silver, Aja Huang, Chris J. Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, et al. Mastering the game of Go with deep neural networks and tree search. *Nature*, 529:484–489, 2016.